

Bio-Entropic Web Data Collection and Market Intelligence Framework

Empirical Verification, A/B Control Experimentation, and Systemic Vulnerability Analysis

LemonScripter Systems Architecture & Advanced Automation Engineering

September 2026

Abstract

Modern behavioral web application firewalls and anti-bot defenses (e.g., Akamai Bot Manager, Cloudflare Turnstile, AWS WAF) effectively detect traditional deterministic automation scripts and uniform pseudo-random number generators (PRNG). This whitepaper presents a production-grade data collection and market intelligence architecture driven by a **Bio-Entropic State Engine**. By ingesting multi-day (168-hour) circadian and ultradian biological actigraphy datasets, the framework dictates both macroscopic operating schedules and microscopic non-deterministic timing jitter. We present verifiable empirical data from live production runs on target e-commerce architectures, provide an automated A/B control experiment demonstrating the mitigation of silent quarantine (shadow-banning), and conduct a rigorous engineering analysis of potential systemic failure modes and their respective countermeasures.

1 System Architecture and Operating Principles

The framework is architected into four strictly decoupled, robust subsystems:

1. **Bio-Entropic Timing & Seeding Engine (`src/bio_engine.py`):** Translates a 168-hour biological actigraphy time-series into macro-operational states, partitioning activity into *HUNTING_EXPLORATION* (active extraction) and *DIGESTING_REST* (sleep/quiescence) windows. Microsecond-level delays are computed via biological log-normal distribution equations calibrated to metabolic velocity rather than linear pseudo-random noise.
2. **Stochastic State Machine & Movement Engine (`src/behavioral_motor.py`):** Drives continuous, sub-pixel cursor movement via cubic and quartic Bézier polynomials governed by Fitts's Law and minimum-jerk acceleration models. It incorporates neuromuscular micro-tremor and stochastic attention lapses, eliminating artificial linear mouse trajectories.
3. **Hardened Stealth & Network Layer (`src/stealth_layer.py`):** Eliminates standard automation telemetry (`navigator.webdriver = undefined`), normalizes hardware traits (`hardwareConcurrency`, `deviceMemory`), and injects imperceptible ± 1 least significant bit (LSB) mathematical perturbations into Canvas, AudioContext, and WebGL rendering pipelines.
4. **Persistence & Data Pipeline (`src/data_pipeline.py`):** Stores relational market intelligence, time-series price histories, and review sentiments into an ACID-compliant SQLite

schema, while preserving cryptographic session tokens and cookies across runs using Playwright `storage_state` mechanics.

2 Empirical Verification and Live Production Metrics

During live verification sweeps conducted on Amazon’s production infrastructure, the framework successfully executed complete extraction cycles without encountering rate limits, blocks, or CAPTCHA challenges:

Metric Description	Recorded Value	Engineering Significance
Unique extracted products	43 ASINs	Unhindered, complete SERP HTML ingestion
Price history time-slices	45 snapshots	Relational time-series persistence (<code>market_data.db</code>)
Mean product price	\$249.38	Market spectrum (\$54.18 – \$903.51)
Mean customer rating	4.47 / 5.0	Authentic customer review data captured
Session state payload	13.5 KB	39 server-issued Amazon cookies authenticated
Micro-jitter latency window	62.8 – 123.4 ms	Non-deterministic log-normal variance

Table 1: Empirical telemetry and storage metrics from live production execution

The captured session state file (`session_state.json`) contains cryptographically signed Amazon tokens (`session-id`, `ubid-main`, `sst-main`), which are only issued upon passing initial TLS and client-side JavaScript behavioral inspection.

3 A/B Control Experimentation

To formally validate anti-detection efficacy, the suite includes an automated control harness (`counter_test.py`) executing identical network routes with both a traditional uncloaked bot and the bio-entropic stealth agent:

Evaluated Parameter	[A] Naive Traditional Bot	[B] Bio-Entropic Stealth Agent
navigator.webdriver WebGL Renderer	True (Exposed) Default Software Headless	None (Clean / Undefined) ANGLE (NVIDIA RTX 3070)
Issued Session Cookies	0 cookies (Quarantined)	38 cookies (Authenticated)
Cursor Movement Physics	0 ms Instantaneous Teleport	Bézier S-Curve + Fitts’s Law
Timing Characteristics	Static Linear Determinism	Biological Log-Normal Jitter
WAF Defense Outcome	Silent Shadow-Banning	100% Clean Pass (No Challenge)

Table 2: Empirical comparison matrix from the A/B control experiment

3.1 The Silent Shadow-Banning Phenomenon

The experiment uncovered a key WAF behavioral mechanic: rather than presenting an immediate visual block or CAPTCHA, the target system silently quarantined the naive automation client by **refusing to issue session and identity cookies (0 cookies recorded)**, returning truncated and degraded responses. In contrast, the Bio-Entropic Agent received 38 fully authenticated cookies and unrestricted catalog access.

4 Systemic Vulnerabilities and Engineering Caveats

To ensure long-term production stability, we analyze four critical attack surfaces and their respective mitigations:

Caveat 1: TLS/TCP JA3/JA4 Protocol Discrepancy

Risk: If the application layer broadcasts a modern Chrome User-Agent while the low-level TLS handshake (Cipher Suite ordering, ALPN, TLS extensions) originates from an unhardened client, edge WAFs flag the connection via passive OS/TLS fingerprinting.

Countermeasure: Utilize Chromium's native network socket implementation or bridge secondary requests through `curl_cffi` / uTLS with matched Chrome fingerprints.

Caveat 2: Residential IP Overutilization & Subnet Reputation

Risk: Sustaining hundreds of requests on a single residential proxy IP without operational pauses triggers subnet-level anomaly scoring.

Countermeasure: Synchronize proxy rotation strictly with biological macro-rest (*DI-GESTING_REST*) intervals, ensuring traffic naturally decays to zero during simulated rest cycles.

Caveat 3: Over-Noised Canvas/WebGL Signatures (Entropy Trap)

Risk: Injecting high-amplitude or deterministic pseudorandom noise into Canvas/WebGL render buffers exposes the client to synthetic fingerprint tampering detection.

Countermeasure: Restrict canvas perturbations strictly to ± 1 least significant bit (LSB) alterations, preserving hardware rendering coherence.

Caveat 4: Cyclical Telemetry Repetition (Machine Learning Profiling)

Risk: Looping an identical 168-hour actigraphy dataset over multiple months enables long-term behavioral classifiers to detect periodic determinism.

Countermeasure: Superimpose continuous stochastic drift (via an Ornstein-Uhlenbeck mean-reverting process) and seasonal phase modulation upon the primary telemetry signal.

5 Conclusion

The Bio-Entropic framework demonstrates that combining empirical biological actigraphy, physical Bézier cursor dynamics, and multi-tier fingerprint hardening establishes a statistically robust methodology for competitive market intelligence, reliably bypassing advanced behavioral detection paradigms.